



## ModelOps RFP Requirements

This document is an example RFP for addressing ModelOps (and MLOps) functional requirements. It is the result of interviews with several industry experts and analysts.

ModelOps (and its MLOps subset which focus on ML models only) is a key capability that is required for successful AI/ML model operations once models have been developed. It is a discipline that is separate and apart from model development. Industry experts and analysts are recognizing that model development and model operations are different disciplines, requiring different capabilities, tools and even teams. Gartner, in a recent article, states, "Platform independence: AI pipelines span multiple environments from developer notebooks to edge to data center to cloud deployments. A true ModelOps framework allows you to bring standardization and scalability across these disparate environments so that development, training and deployment processes can run consistently and in a platform-agnostic manner." Gartner report "Innovation Insight for ModelOps", Farhan Choudhary, Shubhangi Vashisth, Arun Chandrasekaran, Erick Brethenoux. 6 August 2020.

Category	Requirement	Description
<b>Model Management</b> suggested weighting: 20%		
	<b>Model Implementation Agnosticity</b>	Supports the operationalization of all models, inclusive of: <ul style="list-style-type: none"> <li>- All Model Types: AI, ML, mathematical optimization, rules-based, etc.</li> <li>- All Model Languages (and associated Frameworks): Java, Python, R, C/C++, Scala, Matlab, etc.</li> <li>- All Model Execution Environments: Kubernetes/Docker, Java, C/C++, Matlab, SageMaker, Azure ML, GCP, Spark, etc.</li> <li>- All Model Execution Locations: On-Prem, Cloud, Hybrid</li> </ul>
	<b>Standard Model Definition</b>	Provide a consistent definition (and underlying persistence mechanism) of all the core elements that compose a model, regardless of the language/framework, Data Science workbench, underlying infrastructure, or data platform used.
	<b>Custom Metadata</b>	Allow for extension of the core Standard Model definition with custom metadata, which may come from various integrating systems or via user input. The custom metadata must be collected and persisted with each and every version of a model. The custom metadata must be available to be used in a Model Life Cycle to allow an enterprise to define and enforce governance, technical, and business requirements that are specific to their processes.
	<b>Production Model Inventory</b>	Provide a centralized store for viewing, managing, and maintaining all models (post-development) across the enterprise, regardless of the model type, framework, platform, or environment.
	<b>Model Registration</b>	Allow users to onboard their developed model into the ModelOps system by collecting the key elements that compose a model. Model registration should be enabled via CLI, import from a git repository, or via model factory plugins, such as a Jupyter or Rstudio plugin.

<b>Model Artifact Management</b>	Collect and manage all model "artifacts" with the model metadata as part of the "Standard Model Definition" for the registered model. Depending on the type of model that the data scientist uses, the "model training" will result in a set of outputs, which can be coefficients, weights files, binary objects, etc. These are critical parts of the "model" and therefore must be managed in conjunction with the model code. Additionally, the artifact management capability must integrate with common file/binary artifact storage mechanisms such as an Object Store (S3) to provide for scalability and manageability, while also allowing an enterprise to leverage their existing IT investments.
<b>Dependency Management Integration</b>	Manage the list of frameworks, libraries, and other dependencies that are required to execute a given model. While the actual libraries are typically managed in an existing repository (e.g. Artifactory), the ModelOps capability needs to manage--with each version of the model--the exact libraries/frameworks--and the version of the library--for the model. Ensure that the model is deployed to an available execution context that has the required libraries to run the model.
<b>Source Code Management Integration</b>	Integrate with the enterprise's standard for source code management (Bitbucket). Associate details of the model's git assets as metadata with the standard model definition.
<b>Test Results Management</b>	Persist all instances of tests/metrics that have been executed on a per-model basis, regardless of the type of model or the platform upon which it executes. A "record" of a test/metrics job should be tied to an immutable snapshot of a model for history and auditability.
<b>Associated Models</b>	Ability to define associations between models, which may be focused on reusable monitoring "models" (e.g. drift models) across multiple models; or could be associating multiple models in a parent child or linked ensemble approach.

<b>Model Life Cycle</b> suggested weighting: 25%	
<b>Model Life Cycle (MLC) Design</b>	Ability to design and build holistic Model Life Cycles (MLC's) that incorporate business, technical, and compliance/governance requirements across the entire life of a model to allow orchestration and management of the entire AI Governance workflow, from registration to testing, promotion, business approvals, governance approvals, technical approvals, deployment, monitoring, improvement, and eventual retirement. Must provide the ability to create different Model Life Cycles (MLC's) for different groups, classes of models, or even on a per model basis.
<b>Model Life Cycle (MLC) Integration</b>	Ability to integrate with existing operational systems (e.g. Production Support CR/IM ticketing, alerting systems, etc.) and with existing Governance systems (e.g. MRM/Compliance systems). The integration must be systematic (via API's/other) to allow for automation of the overall Model Life Cycle. Must provide the ability to systematically collect information/metadata from the integrating systems and persist as metadata with the specific version of the model.
<b>Model Life Cycle (MLC) Automation</b>	Manage and automate the entire life cycle of a model—from model registration to testing, promotion, reviews, approvals, etc. Integrate with the underlying standard IT systems (ticketing, security scanning, etc.), MRM/compliance systems, and other AI Governance processes/systems to manage and orchestrate the overall AI governance workflow. Additionally, allow for the creation of custom MLCs to cater to the differing needs of each team, business unit, or even types of models.
<b>Model Life Cycle (MLC) Lineage</b>	Collect and persist each step in the model's life cycle for full auditability for each and every version of a model. This includes all of the steps taken to deploy a model to production, including all approvals, promotions, and other operational controls; as well as all steps involved in running on-going production operational controls, monitors, and metrics.

<b>Model Life Cycle (MLC) Reusable Operational Patterns</b>	Built-in customized Model Life Cycle definitions, which are prepackaged automated MLC to shorten the time to business, and allow for reusability across the enterprise for multiple use-cases.
<b>Model Approvals (technical, business and compliance)</b>	Automate the specific approval process(es) required to get a model into business. Integrate with the enterprise's standard change and incident management systems (e.g. ServiceNow) as well as task ticketing systems (e.g. JIRA), and Model Risk Management (MRM)/Compliance systems, to automate and track approvals across all business, data science, operations and IT teams.
<b>Unencoded Thresholds Definitions</b>	Ability to define--via decision tables or other rules--specific thresholds within which a model should operate for various points in its Model Life Cycle. For example, thresholds may be set for matching compliance rules back to statistical performance to provide clear but comprehensive rules for operating the model within the bounds of the compliance and business requirements. Thresholds definitions do not require programming skillset to be decodified.
<b>Code and Quality Assurance Integration</b>	Integrate with enterprise security scanning software (e.g. SonarQube) and/or code scanning (e.g. Veracode) to streamline operationalization of a given model. Additionally, allow for automation of the security scanning process as part of a given Model Life Cycle.
<b>Model Governance</b> suggested weighting: 25%	
<b>Model "Snapshots" ("Versions")</b>	Provide the ability to systematically take a snapshot of a model in time, including all of the model's source code, artifacts, documentation, and other metadata. This snapshot must be immutable and maintained in perpetuity for long-term auditability. Each Production model must have traceability to a specific immutable snapshot. The snapshot--and all of the associate metadata for that model snapshot--can be exported for reporting purposes.

<b>Model Reproducibility</b>	Allow for rapid systematic reproduction of training, evaluation, and/or scoring of a particular model version, especially of models that are currently or previously deployed in Production. The specific version of the relevant code, artifacts, and other model assets must be used to conduct the reproducibility test.
<b>Continuous Compliance Checking - Define Rules</b>	Ability to define--via decision tables--specific conditions in which a model must operate to be within compliance. Must be able to define the rules on a per-model basis which are persisted, maintained, and orchestrated with each and every version of a given model. These rules must be able to incorporate custom metadata for a given model, which may have been pulled in automatically via an MRM/compliance/other governance system.
<b>Continuous Compliance Checking - Enforcing Rules</b>	Automated enforcement of all compliance/governance rules that have been defined on a per model basis. Breach of any particular rule can trigger alerts/notifications as dictated by the Model Life Cycle.
<b>Interpretability Integration</b>	Integrate with leading frameworks (Shap, KLime, etc.) to provide a consistent and automated approach to identifying feature importance for models for a given version of a model.
<b>Ethical Fairness/Bias Integration</b>	Integrate with leading frameworks (Aequitas) to measure and monitor a model's fairness in providing positive outcomes for all protected and/or sensitive classes.
<b>Per model comprehensive audit trail</b>	Support full auditability for each and every version of a model. Capture each step of the entire process to deploy a model to production, including all approvals, promotions, and other operational controls; as well as all steps involved in running on-going production operational controls, monitors, and metrics.

<b>Model Execution &amp; Monitoring</b> suggested weighting: 15%	
<b>Support for multiple model execution platforms</b>	Integrate with multiple model execution (DSML) platforms, providing automated orchestration of jobs, deployment, and monitoring. The integration should occur regardless of the underlying framework (scikitlearn, R, xgboost, TensorFlow, pytorch, etc.) and language (R, python, Scala, C, etc.). The capability should support all existing models--whether they are machine learning, statistical, or even rules-based and regardless of where the model execute: Cloud, on Prem, Hybrid.
<b>Support for multiple modes of scoring: Batch, Request/Response, Streaming</b>	Supports batch, request/response (e.g. REST), and streaming model execution—without changing the underlying algorithm (the “math”). Support switching of scoring modes based on the current stage of model development without changing the model's math: e.g. during early development, a data scientist may want to quickly test using a REST deployment; however, during testing, UAT, and production, they actually need to execute the model in batch due to data volumes.
<b>Job Orchestration &amp; Monitoring</b>	For all models in the ModelOps production model inventory, provides the overarching job orchestratrion and monitoring, regardless of the model execution platform. In particular, a "record" of a job--and any available status, metrics, and results of that job--should be tied to an immutable snapshot of a model for auditability.
<b>Alerting and Automated Notifications</b>	Provide ability to automatically raise alerts and notifications based on model-specific errors, model execution errors--including infrastructure/platform issues, or specific actions/steps in the model life cycle. Typically these alerts would be tied to thresholds (see above) to allow for clear, but comprehensive, orchestration of when alerts are triggered. The ModelOps capability should be flexible to integrate with existing operational alerting and/or production ticketing systems to enable L1 support to provide 24x7 production support for all models across the enterprise.

<b>ModelOps Life Cycle KPI / Process Monitoring</b>	Collect and aggregate detailed metrics about Model Life Cycle execution across the enterprise, including specific times for each step in a Model Life Cycle, allowing for understanding of where there are bottlenecks in the process so that they can address accordingly.
<b>Comprehensive Monitoring</b>	Provides comprehensive monitoring of a model running in Production, including: - Data Drift Monitoring - Model Concept Drift - Statistical Performance Monitoring - Ethical Fairness Drift - Business Simulation (what if analysis) - Production Comparisons (e.g. Parallel Model Execution) Users must be able to provide their own custom metrics (i.e. should not be locked into specific metrics that the ModelOps software allows). Additionally, the user must be able to set thresholds to evaluate the monitoring metrics, including comparisons to baseline data (e.g. compare a current window of production inference requests to the training data input feature distribution. Finally, the system must be capable of triggering automated alerts if thresholds are exceeded.
<b>Custom Metrics Monitoring</b>	Allow a data scientist or ModelOps engineer to define custom statistics to be captured about the model.
<b>Trigger model refresh or retraining based on monitoring alerts</b>	Allow users to monitor specific statistics about a model, and based on pre-defined thresholds, trigger an automated re-training process if those thresholds are exceeded.
<b>Champion/Challenger</b>	Allow for the comparison of the current "champion" production model against a proposed "better" model, using a variety of statistical and technical performance metrics (e.g. AUC, RMSE, F1/F2, ROC Curve, Confusion Matrix).



<b>Visualization</b> suggested weighting: 5%		
	<b>ModelOps Console</b>	Provide a robust ModelOps Console that provides the ability to manage all models, Model Life Cycles, and model runtimes across the enterprise, as well as detailed model metrics such as statistical metrics, drift metrics, and bias metrics.
	<b>Visualization: BI tools integrations</b>	Ability to integrate with BI tools to create custom reports to gain visibility of aggregated and detailed views on model performance against statistical, business, and risk thresholds.
<b>Non-Functional Requirements</b> suggested weighting: 10%		
	<b>Product Architecture</b>	ModelOps capability should be designed to be flexible and extensible, using a modern micro services architecture, to allow enterprises to customize and integrate the ModelOps services into existing systems, platforms, and processes.
	<b>Infrastructure</b>	Ability to run Software anywhere: on Cloud, on Prem, Hybrid.
	<b>Authentication &amp; Group-based access</b>	Group-based access control to isolate which teams can see specific models, model assets, runtimes with specific deployed models, test results, integrating with the enterprise's oauth2 and LDAP/AD systems for authentication and access control. Must tie into the enterprise's SSO capability.
	<b>Systematic Access (API's)</b>	Provide a comprehensive API to allow enterprises to create custom Model Life Cycles or custom integrations into consuming business applications. Expose core ModelOps services as RESTful API's allowing enterprises to easily customize the use of the ModelOps solution to their specific business requirements.
	<b>Ownership</b>	Allow for definition of ownership for model operationalization, by providing a console for 24x7 model management in production, and with visibility into the model path to/in production with pre-defined ownership of responsibilities, triage and remediation processes.
	<b>Localization</b>	Support for the linguistic, technical, and overall cultural requirements for all regions of a global enterprise